

Data Wrangling I

Perform the following operations using Python on any open source dataset (e.g., data.csv)

1. Import all the required Python Libraries.
2. Locate an open source data from the web (e.g., <https://www.kaggle.com>). Provide a clear description of the data and its source (i.e., URL of the web site).
3. Load the Dataset into pandas dataframe.
4. Data Preprocessing: check for missing values in the data using pandas `isnull()`, `describe()` function to get some initial statistics. Provide variable descriptions. Types of variables etc. Check the dimensions of the data frame.
5. Data Formatting and Data Normalization: Summarize the types of variables by checking the data types (i.e., character, numeric, integer, factor, and logical) of the variables in the data set. If variables are not in the correct data type, apply proper type conversions.
6. Turn categorical variables into quantitative variables in Python. In addition to the codes and outputs, explain every operation that you do in the above steps and explain everything that you do to import/read/scrape the data set.

1.1 import all the required Python Libraries

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
```

Locate an open source data from the web (e.g., <https://www.kaggle.com>). Provide a clear description of the data and its source (i.e., URL of the web site).

```
url_link="https://raw.githubusercontent.com/rohinidevkar/DSBDA/main/autodata.csv"
df = pd.read_csv(url_link)
```

```
df.head(10)
```

	Unnamed: 0	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location
0	0	3	122	alfa-romero	std	two	convertible	rwd	
1	1	3	122	alfa-romero	std	two	convertible	rwd	
2	2	1	122	alfa-romero	std	two	hatchback	rwd	
3	3	2	164	audi	std	four	sedan	fwd	
4	4	2	164	audi	std	four	sedan	4wd	

df.tail()

	Unnamed: 0	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location
196	196	-1	95	volvo	std	four	sedan	rwd	front-engine
197	197	-1	95	volvo	turbo	four	sedan	rwd	front-engine
198	198	-1	95	volvo	std	four	sedan	rwd	front-engine
199	199	-1	95	volvo	turbo	four	sedan	rwd	front-engine
200	200	-1	95	volvo	turbo	four	sedan	rwd	front-engine

5 rows × 30 columns

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 201 entries, 0 to 200
Data columns (total 30 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            201 non-null   int64
1   symboling              201 non-null   int64
2   normalized-losses     201 non-null   int64
3   make                   201 non-null   object
4   aspiration              201 non-null   object
5   num-of-doors           201 non-null   object
6   body-style             201 non-null   object
7   drive-wheels           201 non-null   object
8   engine-location        201 non-null   object
9   wheel-base             201 non-null   float64
10  length                 201 non-null   float64
11  width                  201 non-null   float64
12  height                 201 non-null   float64
13  curb-weight            201 non-null   int64
14  engine-type            201 non-null   object
15  num-of-cylinders       201 non-null   object
```

```

16 engine-size      201 non-null    int64
17 fuel-system      201 non-null    object
18 bore             201 non-null    float64
19 stroke           197 non-null    float64
20 compression-ratio 201 non-null    float64
21 horsepower       199 non-null    float64
22 peak-rpm         199 non-null    float64
23 city-mpg         201 non-null    int64
24 highway-mpg      201 non-null    int64
25 price            201 non-null    float64
26 city-L/100km     201 non-null    float64
27 horsepower-binned 199 non-null    object
28 diesel           201 non-null    int64
29 gas              201 non-null    int64
dtypes: float64(11), int64(9), object(10)
memory usage: 47.2+ KB

```

```
df.describe()
```

	Unnamed: 0	symboling	normalized-losses	wheel-base	length	width	height
count	201.000000	201.000000	201.00000	201.000000	201.000000	201.000000	201.000
mean	100.000000	0.840796	122.00000	98.797015	0.837102	0.915126	53.766
std	58.167861	1.254802	31.99625	6.066366	0.059213	0.029187	2.447
min	0.000000	-2.000000	65.00000	86.600000	0.678039	0.837500	47.800
25%	50.000000	0.000000	101.00000	94.500000	0.801538	0.890278	52.000
50%	100.000000	1.000000	122.00000	97.000000	0.832292	0.909722	54.100
75%	150.000000	2.000000	137.00000	102.400000	0.881788	0.925000	55.500
max	200.000000	3.000000	256.00000	120.900000	1.000000	1.000000	59.800

```
df.isnull()
```

	Unnamed: 0	symboling	normalized- losses	make	aspiration	num- of- doors	body- style	drive- wheels	engi locat
0	False	False	False	False	False	False	False	False	Fa
1	False	False	False	False	False	False	False	False	Fa
2	False	False	False	False	False	False	False	False	Fa

```
df.isnull().sum()
```

Unnamed: 0	0
symboling	0
normalized-losses	0
make	0
aspiration	0
num-of-doors	0
body-style	0
drive-wheels	0
engine-location	0
wheel-base	0
length	0
width	0
height	0
curb-weight	0
engine-type	0
num-of-cylinders	0
engine-size	0
fuel-system	0
bore	0
stroke	4
compression-ratio	0
horsepower	2
peak-rpm	2
city-mpg	0
highway-mpg	0
price	0
city-L/100km	0
horsepower-binned	2
diesel	0
gas	0
dtype: int64	

```
df.notnull()
```

	Unnamed: 0	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location
0	True	True	True	True	True	True	True	True	T
1	True	True	True	True	True	True	True	True	T
2	True	True	True	True	True	True	True	True	T
3	True	True	True	True	True	True	True	True	T
4	True	True	True	True	True	True	True	True	T
...

```
df.notnull().sum()
```

```

Unnamed: 0      201
symboling      201
normalized-losses 201
make           201
aspiration     201
num-of-doors   201
body-style     201
drive-wheels   201
engine-location 201
wheel-base    201
length        201
width          201
height        201
curb-weight    201
engine-type    201
num-of-cylinders 201
engine-size    201
fuel-system    201
bore           201
stroke        197
compression-ratio 201
horsepower     199
peak-rpm       199
city-mpg       201
highway-mpg    201
price          201
city-L/100km   201
horsepower-binned 199
diesel         201
gas            201
dtype: int64

```

```

#calculate the mean value for "stroke" column
avg_stroke = df["stroke"].astype("float").mean(axis = 0)
print("Average of stroke :",avg_stroke)

```

```

#replace NaN by mean value in "stroke" column
df["stroke"].replace(np.nan, avg_stroke,inplace = True)

```

```
Average of stroke : 3.2569035532994857
```

Calculate the mean value for the 'horsepower' column :

```
avg_hp=df["horsepower"].astype("float").mean(axis = 0)
print("Average of stroke :",avg_hp)
```

Average of stroke : 103.39698492462311

```
df['horsepower'].replace(np.nan,avg_hp,inplace = True)
```

```
from contextlib import nullcontext
df['num-of-doors'].value_counts()
```

```
four      115
two        86
Name: num-of-doors, dtype: int64
```

```
df['num-of-doors'].value_counts().idxmax()
```

'four'

```
# replace the missing 'num-of-door' values by most frequent
df['num-of-doors'].replace(np.nan, "four" , inplace=True)
```

```
#simply drop whole row with nan in "Horsepower-binned" column
df.dropna(subset=['horsepower-binned'], axis=0 , inplace=True)
```

```
#reset index, because we dropped two rows
df.reset_index(drop=True, inplace=True)
```

```
df.isnull().sum()
```

```
Unnamed: 0      0
symboling       0
normalized-losses 0
make           0
aspiration      0
num-of-doors    0
body-style      0
drive-wheels    0
engine-location 0
wheel-base     0
length         0
width          0
height         0
curb-weight     0
engine-type     0
num-of-cylinders 0
engine-size     0
fuel-system     0
bore           0
stroke         0
compression-ratio 0
```

```

horsepower      0
peak-rpm        2
city-mpg        0
highway-mpg     0
price           0
city-L/100km    0
horsepower-binned 2
diesel          0
gas             0
dtype: int64

```

DATA STANDARDIZATION : It is process of transforming data into common format which allows the researcher to make meaningful comparison

```

df['city-L/100km']=235/df['city-mpg']
df.head()

```

	Unnamed: 0	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	drive-wheels	e
0	0	3	122	alfa-romero	std	two	convertible	rwd	
1	1	3	122	alfa-romero	std	two	convertible	rwd	
2	2	1	122	alfa-romero	std	two	hatchback	rwd	
3	3	2	164	audi	std	four	sedan	fwd	
4	4	2	164	audi	std	four	sedan	4wd	

5 rows × 10 columns

```

df['highway-L/100km']=235/df["highway-mpg"]
df.head()

```

Unnamed: 0	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location
0	0	3	122	alfa-romero	std	two	convertible	rwd
1	0	3	122	alfa-romero	std	two	convertible	rwd

DATA NORMALIZATION : It is process of transforming several values into similar range

2	2	1	122	alfa-romero	std	two	hatchback	rwd
---	---	---	-----	-------------	-----	-----	-----------	-----

```
df['length']=df['length']/df['length'].max()
df['width']=df['width']/df['width'].max()
df['height']=df['height']/df['height'].max()
```

```
df[['length','width','height']].head()
```

	length	width	height
0	0.811148	0.890278	0.816054
1	0.811148	0.890278	0.816054
2	0.822681	0.909722	0.876254
3	0.848630	0.919444	0.908027
4	0.848630	0.922222	0.908027

INDIATOR VARIABLE : Indicator variable or dummy variable are used to label numerical variable used to label categories

```
df.columns
```

```
Index(['Unnamed: 0', 'symboling', 'normalized-losses', 'make', 'aspiration',
      'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',
      'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',
      'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
      'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
      'highway-mpg', 'price', 'city-L/100km', 'horsepower-binned', 'diesel',
      'gas', 'highway-L/100km'],
      dtype='object')
```

```
df['aspiration'].value_counts()
```

```
std      163
turbo     36
Name: aspiration, dtype: int64
```

```
dummy_var_1=pd.get_dummies(df['aspiration'])
dummy_var_1.head()
```


	std	turbo
0	1	0
1	1	0
2	1	0
3	1	0
4	1	0

```
df=pd.concat([df,dummy_var_1], axis=1)
df.drop('aspiration',axis = 1 , inplace = True)
```

```
df.head()
```

	Unnamed: 0	symboling	normalized-losses	make	num-of-doors	body-style	drive-wheels	engine-location	wheel-base
0	0	3	122	alfa-romero	two	convertible	rwd	front	88.5
1	1	3	122	alfa-romero	two	convertible	rwd	front	88.5
2	2	1	122	alfa-romero	two	hatchback	rwd	front	94.5
3	3	2	164	audi	four	sedan	fwd	front	95.5
4	4	2	164	audi	four	sedan	4wd	front	95.5

5 rows × 52 columns

The last columns are indicator variable which are represented by 0's and 1's

```
df.columns
```

```
Index(['Unnamed: 0', 'symboling', 'normalized-losses', 'make', 'num-of-doors',
      'body-style', 'drive-wheels', 'engine-location', 'wheel-base', 'length',
      'width', 'height', 'curb-weight', 'engine-type', 'num-of-cylinders',
      'engine-size', 'fuel-system', 'bore', 'stroke', 'compression-ratio',
      'horsepower', 'peak-rpm', 'city-mpg', 'highway-mpg', 'price',
      'city-L/100km', 'horsepower-binned', 'diesel', 'gas', 'highway-L/100km',
      'std', 'turbo', 'std', 'turbo', 'std', 'turbo', 'std', 'turbo', 'std',
      'turbo', 'std', 'turbo', 'std', 'turbo', 'std', 'turbo', 'std', 'turbo',
      'std', 'turbo', 'std', 'turbo'],
      dtype='object')
```


BINNING: It is process of transforming continous data into discrete categorical 'bins' for group analysis

```
df ["horsepower"]=df ["horsepower"].astype(float, copy=True)
```

```
%matplotlib inline
import matplotlib.pyplot as plt
from matplotlib import pyplot
import numpy as np

plt.matplotlib.pyplot.hist(df['wheel-base'])
plt.matplotlib.pyplot.xlabel('wheel-base')
plt.matplotlib.pyplot.ylabel('count')
plt.matplotlib.pyplot.title('wheel-base bins')
```

```
Text(0.5, 1.0, 'wheel-base bins')
```



```
bins = np.linspace(min(df['wheel-base']),max(df['wheel-base']),4)
bins
```

```
array([ 86.6          ,  98.03333333, 109.46666667, 120.9          ])
```

