

## ▼ Data Analytics I

Create a Linear Regression Model using Python/R to predict home prices using Boston Housing Dataset (<https://www.kaggle.com/c/boston-housing>). The Boston Housing dataset contains information about various houses in Boston through different parameters. There are 506 samples and 14 feature variables in this dataset.

The objective is to predict the value of prices of the house using the given features

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
from sklearn.datasets import load_boston
boston = load_boston()
```

/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning

The Boston housing prices dataset has an ethical problem. You can refer to the documentation of this function for further details.

The scikit-learn maintainers therefore strongly discourage the use of this dataset unless the purpose of the code is to study and educate about ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original source::

```
import pandas as pd
```

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu ✕

```
data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
data = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]
```

Alternative datasets include the California housing dataset (i.e. :func:`~sklearn.datasets.fetch\_california\_housing`) and the Ames housing dataset. You can load the datasets as follows::

```
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()
```

for the California housing dataset and::

```
from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)
```

for the Ames housing dataset.

```
warnings.warn(msg, category=FutureWarning)
```

```
boston.data.shape
```

(506, 13)

```
boston.feature_names
```

```
array(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD',  
      'TAX', 'PTRATIO', 'B', 'LSTAT'], dtype='<U7')
```

```
data = pd.DataFrame(boston.data)  
data.columns = boston.feature_names
```

```
data.head(15)
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90
5	0.02985	0.0	2.18	0.0	0.458	6.430	58.7	6.0622	3.0	222.0	18.7	394.12
6	0.08829	12.5	7.87	0.0	0.524	6.012	66.6	5.5605	5.0	311.0	15.2	395.60
7	0.14455	12.5	7.87	0.0	0.524	6.172	96.1	5.9505	5.0	311.0	15.2	396.90
8	0.17801	12.5	7.87	0.0	0.524	6.359	92.9	6.0621	5.0	311.0	15.2	386.63
9	0.17801	12.5	7.87	0.0	0.524	6.359	92.9	6.0621	5.0	311.0	15.2	386.71
10	0.22489	12.5	7.87	0.0	0.524	6.377	94.3	6.3467	5.0	311.0	15.2	392.52
11	0.11747	12.5	7.87	0.0	0.524	6.009	82.9	6.2267	5.0	311.0	15.2	396.90
12	0.09378	12.5	7.87	0.0	0.524	5.889	39.0	5.4509	5.0	311.0	15.2	390.50
13	0.62976	0.0	8.14	0.0	0.538	5.949	61.8	4.7075	4.0	307.0	21.0	396.90
14	0.63796	0.0	8.14	0.0	0.538	6.096	84.5	4.4619	4.0	307.0	21.0	380.02

```
boston.target.shape
```

(506,)

```
data['Price'] = boston.target  
data.head()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	L
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	

```
data.describe()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	L
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574900	4.700050	1.318710	306.858214	18.454381	393.280558	393.310016
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	1.612150	0.832843	81.613214	7.072033	56.341381	57.145549
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.600000	1.000000	222.000000	12.100000	318.100000	312.099996
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	4.000000	1.000000	242.000000	15.000000	396.900000	392.830000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	4.690000	1.000000	296.000000	17.000000	396.900000	392.830000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.000000	2.000000	318.100000	18.000000	396.900000	392.830000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	6.062200	3.000000	318.100000	18.700000	396.900000	392.830000

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 506 entries, 0 to 505
```

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu

```
0  CRIM    506 non-null    float64  
1  ZN      506 non-null    float64  
2  INDUS   506 non-null    float64  
3  CHAS    506 non-null    float64  
4  NOX     506 non-null    float64  
5  RM      506 non-null    float64  
6  AGE     506 non-null    float64  
7  DIS     506 non-null    float64  
8  RAD     506 non-null    float64  
9  TAX     506 non-null    float64  
10 PTRATIO 506 non-null    float64  
11 B       506 non-null    float64  
12 LSTAT   506 non-null    float64  
13 Price   506 non-null    float64  
dtypes: float64(14)  
memory usage: 55.5 KB
```

```
x=boston.data  
y=boston.target  
  
from sklearn.model_selection import train_test_split
```

```
xtrain,xtest,ytrain,ytest = train_test_split(x,y, test_size = 0.2)
```


```
print("xtrain shape :",xtrain.shape)
print("xtest shape :",xtest.shape)
print("ytrain shape :",ytrain.shape)
print("ytest shape :", ytest.shape)
```

```
xtrain shape : (404, 13)
xtest shape : (102, 13)
ytrain shape : (404,)
ytest shape : (102,)
```

```
from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(xtrain,ytrain)
```

```
y_pred = regressor.predict(xtest)
```

```
plt.scatter(ytest,y_pred, c = 'orange')
plt.xlabel("Price : in $1000's")
plt.ylabel("Predicted value")
plt.title("True value vs Predicted value : Linear Regression")
plt.show()
```



```
from sklearn.metrics import mean_squared_error
mse = mean_squared_error(ytest,y_pred)
print("Mean Square Error :",mse)
```

➡ Mean Square Error : 28.89318492981834

✓ 0s completed at 2:27 PM ● ✕

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu ✕