

## GROUP B

### EXPERIMENT : 11

**AIM:** Write a code in JAVA for a simple WordCount application that counts the number of occurrences of each word in a given input set using the Hadoop MapReduce framework on local-standalone set-up.

#### PRE-REQUISITE:

- **Java Installation** - Check whether the Java is installed or not using the following command.  
`java -version`
- **Hadoop Installation** - Check whether the Hadoop is installed or not using the following  
`hadoop version`

#### THEORY:

Steps to execute MapReduce word count :

- Create a text file in your local machine and write some text into it.  
`$ nano data.txt`
- Check the text written in the data.txt file.  
`$ cat data.txt`

File: WC\_Mapper.java

1. `package com.javatpoint;`
2. `import java.io.IOException;`
3. `import java.util.StringTokenizer;`
4. `import org.apache.hadoop.io.IntWritable;`
5. `import org.apache.hadoop.io.LongWritable;`
6. `import org.apache.hadoop.io.Text;`
7. `import org.apache.hadoop.mapred.MapReduceBase;`
8. `import org.apache.hadoop.mapred.Mapper;`
9. `import org.apache.hadoop.mapred.OutputCollector;`
10. `import org.apache.hadoop.mapred.Reporter;`

11.

```
public class WC_Mapper extends MapReduceBase implements Mapper<LongWritable,Text,
Text,IntWritable>{
```

12. private final static IntWritable one = new IntWritable(1);

13. private Text word = new Text();

14.

```
public void map(LongWritable key, Text value,OutputCollector<Text,IntWritable>
output,
```

15. Reporter reporter) throws IOException{

16. String line = value.toString();

17. StringTokenizer tokenizer = new StringTokenizer(line);

18. while (tokenizer.hasMoreTokens()){

19. word.set(tokenizer.nextToken());

20. output.collect(word, one);

21. }

22. }

23. }

File: WC\_Reducer.java

1. package com.javatpoint;

2. import java.io.IOException;

3. import java.util.Iterator;

4. import org.apache.hadoop.io.IntWritable;

5. import org.apache.hadoop.io.Text;

6. import org.apache.hadoop.mapred.MapReduceBase;

7. import org.apache.hadoop.mapred.OutputCollector;

8. import org.apache.hadoop.mapred.Reducer;

9. import org.apache.hadoop.mapred.Reporter;

10.

11.

```
public class WC_Reducer extends MapReduceBase implements Reducer<Text,Int
Writable,Text,IntWritable> {
```

12.

```
    public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text, IntWritable> output,
```

13. Reporter reporter) throws IOException {

14. int sum=0;

15. while (values.hasNext()) {

16. sum+=values.next().get();

17. }

18. output.collect(key, new IntWritable(sum));

19. }

20. }

File: WC\_Runner.java

1. package com.javatpoint;

2.

3. import java.io.IOException;

4. import org.apache.hadoop.fs.Path;

5. import org.apache.hadoop.io.IntWritable;

6. import org.apache.hadoop.io.Text;

7. import org.apache.hadoop.mapred.FileInputFormat;

8. import org.apache.hadoop.mapred.FileOutputFormat;

9. import org.apache.hadoop.mapred.JobClient;

10. import org.apache.hadoop.mapred.JobConf;

11. import org.apache.hadoop.mapred.TextInputFormat;

12. import org.apache.hadoop.mapred.TextOutputFormat;

13. public class WC\_Runner {

14. public static void main(String[] args) throws IOException{

15. JobConf conf = new JobConf(WC\_Runner.class);

16. conf.setJobName("WordCount");

17. conf.setOutputKeyClass(Text.class);

18. conf.setOutputValueClass(IntWritable.class);

19. conf.setMapperClass(WC\_Mapper.class);

20. conf.setCombinerClass(WC\_Reducer.class);

21. conf.setReducerClass(WC\_Reducer.class);

22. conf.setInputFormat(TextInputFormat.class);

```

23.     conf.setOutputFormat(TextOutputFormat.class);
24.     FileInputFormat.setInputPaths(conf,new Path(args[0]));
25.     FileOutputFormat.setOutputPath(conf,new Path(args[1]));
26.     JobClient.runJob(conf);
27. }
28. }

```

## RESULTS:

```

admin1@admin1:~$ mkdir bm
admin1@admin1:~$ cd bm
admin1@admin1:~/bm$ mkdir words
admin1@admin1:~/bm$ javac -classpath hadoop-core-1.2.1.jar words/WordCount.java
admin1@admin1:~/bm$ jar -cvf words.jar -C words/ .
added manifest
adding: WordCount$Reduce.class(in = 1627) (out= 685)(deflated 57%)
adding: WordCount.class(in = 1443) (out= 730)(deflated 49%)
adding: WordCount.java(in = 2055) (out= 657)(deflated 68%)
adding: WordCount$Map.class(in = 1855) (out= 780)(deflated 57%)
admin1@admin1:~/bm$ cd $HADOOP_HOME/sbin
admin1@admin1:~/hadoop-2.7.3/sbin$ ./start-all.sh
This script is deprecated. Instead use start-dfs.sh and start-yarn.sh
Starting namenodes on [localhost]
admin1@localhost's password:
localhost: namenode running as process 6699. Stop it first.
admin1@localhost's password:
localhost: datanode running as process 6893. Stop it first.
Starting secondary namenodes [0.0.0.0]
admin1@0.0.0.0's password:
0.0.0.0: secondarynamenode running as process 7137. Stop it first.
starting yarn daemons
resourcemanager running as process 7318. Stop it first.
admin1@localhost's password:
localhost: nodemanager running as process 7674. Stop it first.
admin1@admin1:~/hadoop-2.7.3/sbin$ hadoop fs -mkdir /input
mkdir: '/input': File exists
admin1@admin1:~/hadoop-2.7.3/sbin$ cd ..
admin1@admin1:~/hadoop-2.7.3$ hadoop fs -put fruits.txt /input
put: 'fruits.txt': No such file or directory
admin1@admin1:~/hadoop-2.7.3$ cd ..
admin1@admin1:~$ hadoop fs -put fruits.txt /input

```

BM steps.pdf

```
localhost: hadoop running as process 6893. Stop it first.
Starting secondary namenodes [0.0.0.0]
admin1@0.0.0.0's password:
0.0.0.0: secondarynamenode running as process 7137. Stop it first.
starting yarn daemons
resourcemanager running as process 7318. Stop it first.
admin1@localhost's password:
localhost: nodemanager running as process 7674. Stop it first.
admin1@admin1:~$hadoop-2.7.3/sbin$ hadoop fs -mkdir /input
mkdir: '/input': File exists
admin1@admin1:~$hadoop-2.7.3/sbin$ cd ..
admin1@admin1:~$hadoop-2.7.3$ hadoop fs -put fruits.txt /input
put: 'fruits.txt': No such file or directory
admin1@admin1:~$hadoop-2.7.3$ cd ..
admin1@admin1:~$ hadoop fs -put fruits.txt /input
put: '/input/fruits.txt': File exists
admin1@admin1:~$ hadoop jar words.jar WordCount /input /output
22/03/29 09:11:04 INFO Configuration.deprecation: session.id is deprecated. Instead, use
dfs.metrics.session-id
22/03/29 09:11:04 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker,
sessionId=
Exception in thread "main" org.apache.hadoop.mapred.FileAlreadyExistsException: Output
directory hdfs://localhost:9000/output already exists
    at
    org.apache.hadoop.mapreduce.lib.output.FileOutputFormat.checkOutputSpecs(FileOutputFormat.ja
va:146)
    at org.apache.hadoop.mapreduce.JobSubmitter.checkSpecs(JobSubmitter.java:266)
    at org.apache.hadoop.mapreduce.JobSubmitter.submitJobInternal(JobSubmitter.java:139)
    at org.apache.hadoop.mapreduce.Job$10.run(Job.java:1290)
    at org.apache.hadoop.mapreduce.Job$10.run(Job.java:1287)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:422)
    at
    org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1698)
    at org.apache.hadoop.mapreduce.Job.submit(Job.java:1287)
```

BM steps.pdf

```
    at org.apache.hadoop.mapreduce.Job.waitForCompletion(Job.java:1308)
    at WordCount.main(WordCount.java:59)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
    at
    sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:498)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:221)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
admin1@admin1:~$ hadoop fs -ls /output
Found 2 items
-rw-r--r-- 1 admin1 supergroup 0 2019-10-31 14:28 /output/_SUCCESS
-rw-r--r-- 1 admin1 supergroup 87 2019-10-31 14:28 /output/part-r-00000
admin1@admin1:~$ hadoop fs -cat /output/part-r-00000
apple 4
banana 2
grapes 3
mango 2
orange 2
pineapple 1
plum 2
pomegranate 1
raspberry 2
```

A terminal window titled 'codegyani@ubuntu64server: ~' with standard window controls. The terminal shows the command 'codegyani@ubuntu64server:~\$ hdfs dfs -cat /r\_output/part-00000' and its output: 'HDFS 1', 'Hadoop 2', 'MapReduce 1', 'a 2', 'is 2', 'of 2', 'processing 1', 'storage 1', 'tool 1', 'unit 1'. The prompt 'codegyani@ubuntu64server:~\$' is followed by a green cursor.

```
codegyani@ubuntu64server:~$ hdfs dfs -cat /r_output/part-00000
HDFS 1
Hadoop 2
MapReduce 1
a 2
is 2
of 2
processing 1
storage 1
tool 1
unit 1
codegyani@ubuntu64server:~$
```

**CONCLUSION:** Hence we have successfully completed hadoop installation and word count program .

## EXPERIMENT : 12

**AIM:** Design a distributed application using MapReduce which processes a log file of a system.

**OBJECTIVES:**

**THEORY:**

**CODE:**

SalesMapper.java

```
package SalesCountry;

import java.io.IOException;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;

public class SalesMapper extends MapReduceBase implements
Mapper<LongWritable, Text, Text, IntWritable> {

    private final static IntWritable one = new IntWritable(1);

    public void map(LongWritable key, Text value,
OutputCollector<Text, IntWritable> output, Reporter reporter) throws
IOException {

        String valueString = value.toString();
        String[] SingleCountryData = valueString.split("-");
        output.collect(new Text(SingleCountryData[0]), one);
    }
}
```

## SalesCountryReducer.java

```
package SalesCountry;

import java.io.IOException;
import java.util.*;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;

public class SalesCountryReducer extends MapReduceBase implements
Reducer<Text, IntWritable, Text, IntWritable> {

    public void reduce(Text t_key, Iterator<IntWritable> values,
OutputCollector<Text,IntWritable> output, Reporter reporter) throws
IOException {

        Text key = t_key;
        int frequencyForCountry = 0;
        while (values.hasNext()) {
            // replace type of value with the actual type of our
value
            IntWritable value = (IntWritable) values.next();
            frequencyForCountry += value.get();

        }
        output.collect(key, new IntWritable(frequencyForCountry));
    }
}
```



## SalesCountryDriver.java

```
package SalesCountry;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;

public class SalesCountryDriver {
    public static void main(String[] args) {
        JobClient my_client = new JobClient();
        // Create a configuration object for the job
        JobConf job_conf = new JobConf(SalesCountryDriver.class);

        // Set a name of the Job
        job_conf.setJobName("SalePerCountry");

        // Specify data type of output key and value
        job_conf.setOutputKeyClass(Text.class);
```

```

        job_conf.setOutputValueClass(IntWritable.class);

        // Specify names of Mapper and Reducer Class
        job_conf.setMapperClass(SalesCountry.SalesMapper.class);

        job_conf.setReducerClass(SalesCountry.SalesCountryReducer.class);

        // Specify formats of the data type of Input and output
        job_conf.setInputFormat(TextInputFormat.class);
        job_conf.setOutputFormat(TextOutputFormat.class);

        // Set input and output directories using command line
arguments,
        //arg[0] = name of input directory on HDFS, and arg[1] =
name of output directory to be created to store the output file.

        FileInputFormat.setInputPaths(job_conf, new Path(args[0]));
        FileOutputFormat.setOutputPath(job_conf, new Path(args[1]));
        my_client.setConf(job_conf);
        try {
            // Run the job
            JobClient.runJob(job_conf);
        } catch (Exception e) {
            e.printStackTrace();
        }
    }
}

```

**OUTPUT:**

```

File Input Format Counters
  Bytes Read=162647
File Output Format Counters
  Bytes Written=3838
hduser@yogesh-X556UOK:~/analyzeLogs$ $HADOOP_HOME/bin/hdfs dfs -cat /output2000/part-00000
18/01/07 14:09:15 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
10.1.1.236 7
10.1.181.142 14
10.1.232.31 5
10.10.55.142 14
10.102.101.66 1
10.103.184.104 1
10.103.190.81 53
10.103.63.29 1
10.104.73.51 1
10.105.160.183 1
10.108.91.151 1
10.109.21.76 1
10.11.131.40 1
10.111.71.20 8
10.112.227.184 6
10.114.74.30 1
10.115.118.78 1
10.117.224.230 1
10.117.76.22 12
10.118.19.97 1
10.118.250.30 7
10.119.117.132 23
10.119.33.245 1
10.119.74.120 1

```

## CONCLUSION:

Hence ,we have successfully performed practical of distributed application using MapReduce